

Overview of CU Research Computing Resources

Nick Featherstone

CU Applied Math

Outline

- Overview
- Storage systems
- Compute resources
- Logging in
- Software stack / Compiling programs
- Submitting batch jobs
- Globus file transfer

What is Research Computing?

- University-wide resource freely available to students/faculty
- High Performance Computing (HPC) Resources for CU
 - Large pool of networked compute nodes
 - (node = motherboard with multiple processors/cores onboard)
 - Parallel mass file storage systems
 - HPC software stack:
 - Commonly used applications: e.g., Matlab, Tensorflow, OpenFOAM
 - Popular Compilers (Intel, GNU, PGI)
 - Optimized, compiled libraries (BLAS, MPI, etc.)
- User support
 - Basic support (how do I log in / use Linux)?
 - Advanced support (code optimization, parallelization)
 - Software installs upon request
 - Instructional seminars (<https://www.colorado.edu/rc/userservices/training>)

Getting Help

- Research Computing Web Page:
 - <https://www.colorado.edu/rc>
- Read the documentation (first!):
 - <https://curc.readthedocs.io/en/latest/>
- Email:
 - rc-help@colorado.edu
- Note:
 - Remember this is a small team of about **10 people** responsible for serving the **entire university's** HPC needs. Sometimes the response time is long. Sometimes you have to iterate before your problem is truly understood.
 - Be cool!

Compute Systems Overview

- Summit

- 10,848 Intel Haswell Cores (4.8 GB/core, 24-cores/node)
- 480 Intel Skylake Cores (7.8 GB/core, 24-core/node)
- And more: <https://www.colorado.edu/rc/articles/rmaccsummit>
- Free-to-use by entire university
- General, low-priority compute 'allocation' provided to every user
- High-priority allocations available to PI's (speak to your advisor)

- Blanca

- Dedicated, grant- or department-funded compute system.
- Heterogeneous node types (Westmere through Skylake)
- Speak to your advisor
- Applied Math has 2, 32-core Intel Skylake nodes
 - Speak to/email Dominique Ingoglia for access: am_itsup@colorado.edu

File Systems Overview

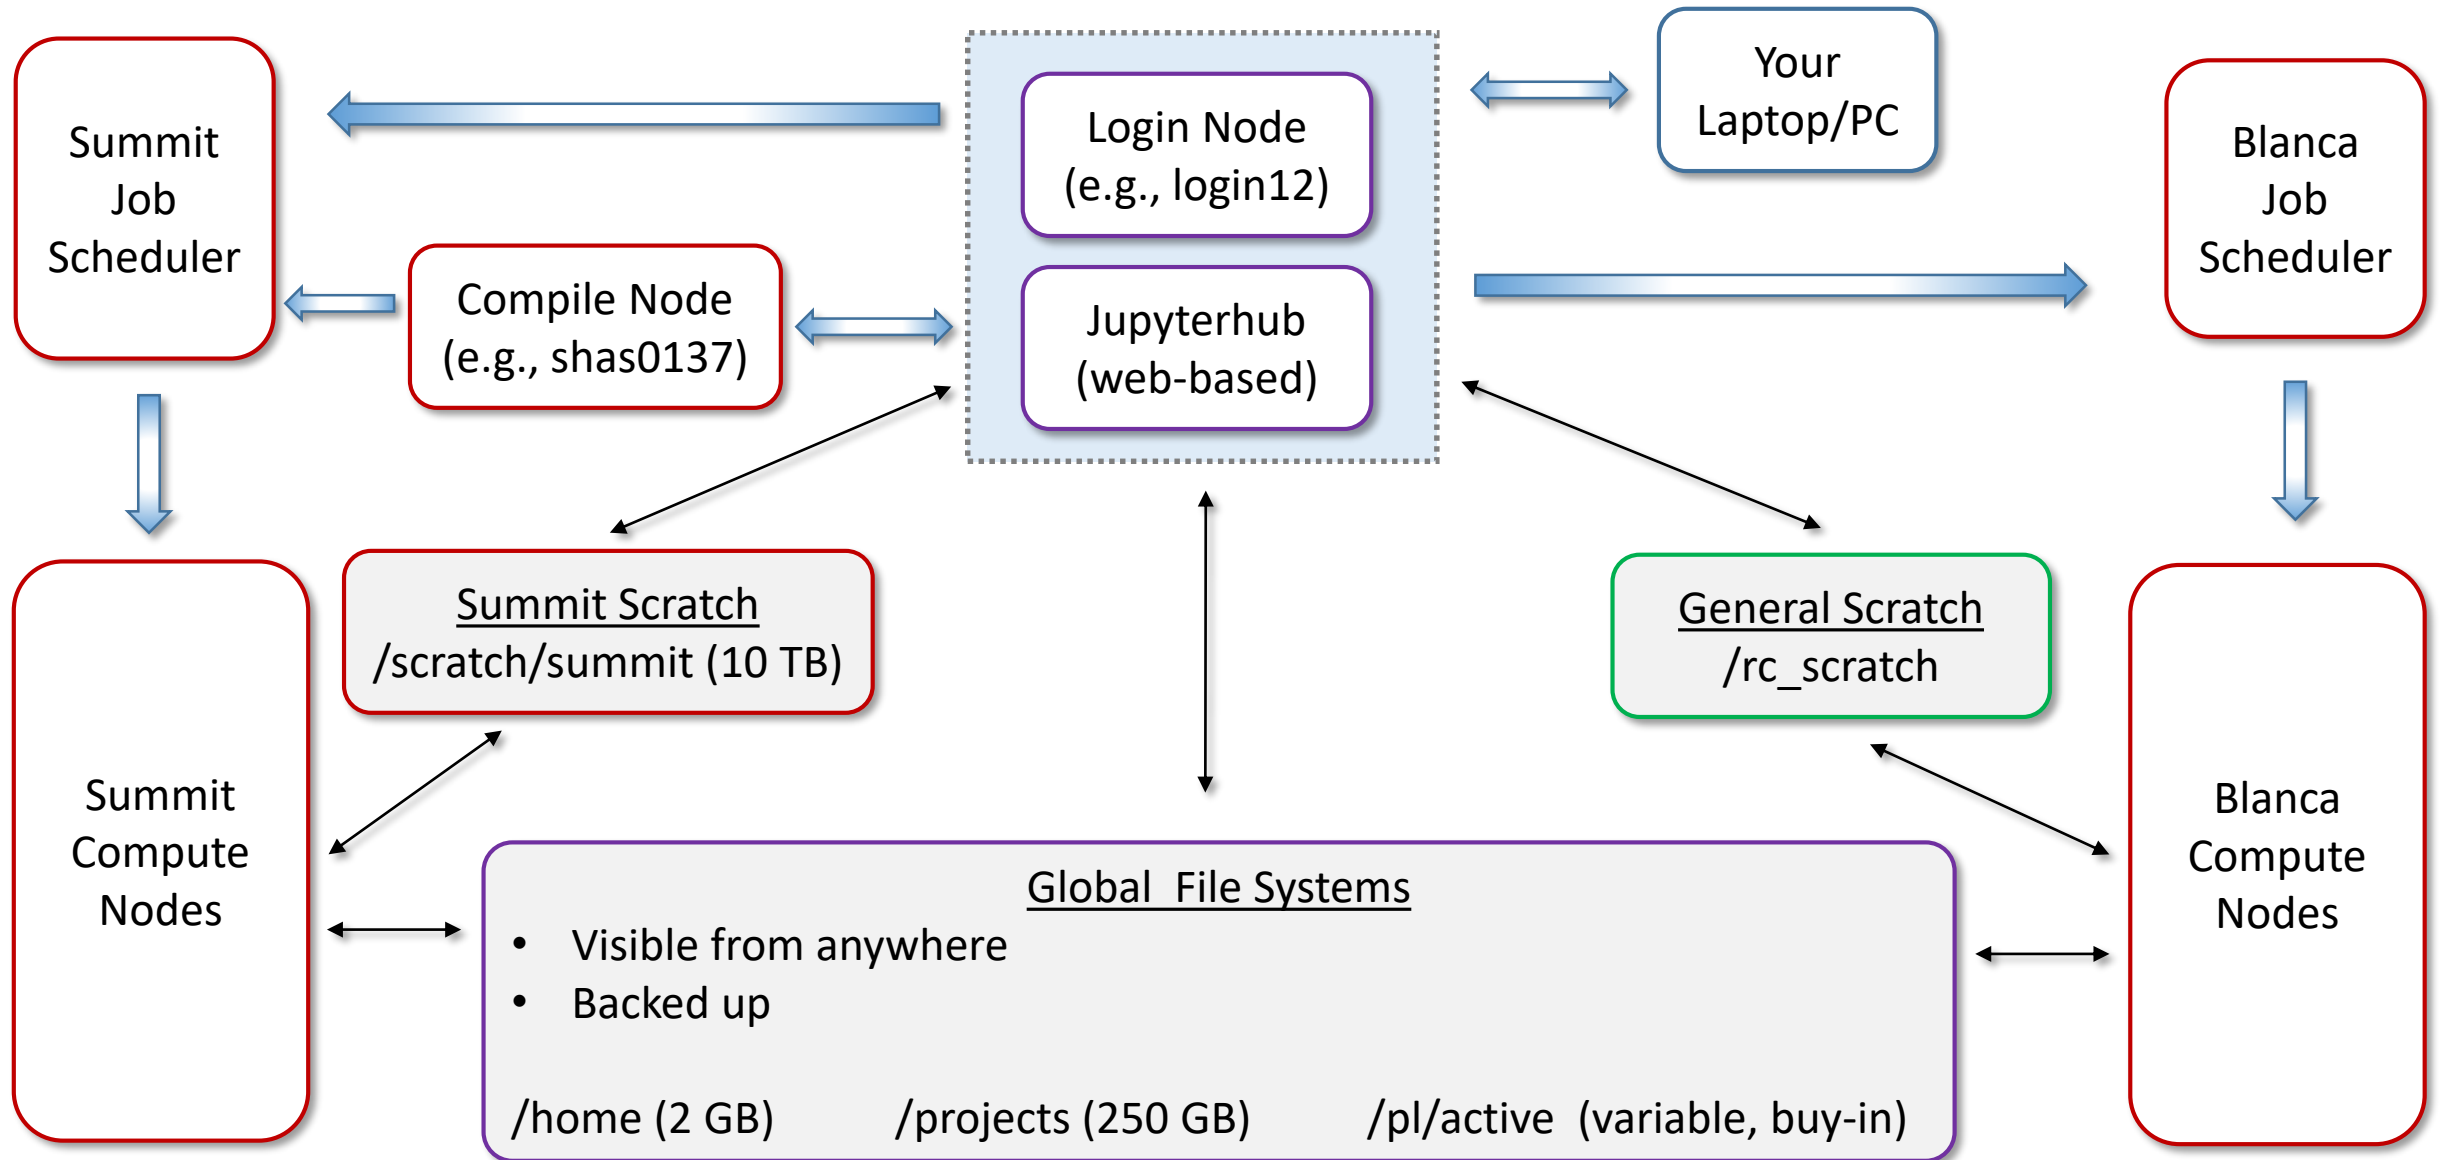
Backed-up Filesystems

- Use these as you see fit, but space is limited
- Non-parallel
- DO NOT have your programs write to these F/S.
- /home/username
 - 2 GB (recommend store source code)
 - Snapshots every 2 hours
- /projects/username
 - 250 GB (permanent input files, etc.)
 - Snapshots every 6 hours
- /pl/active/project_name
 - Ask your advisor (grant or dept. funded)
 - Snapshot frequency varies

Purged Filesystems

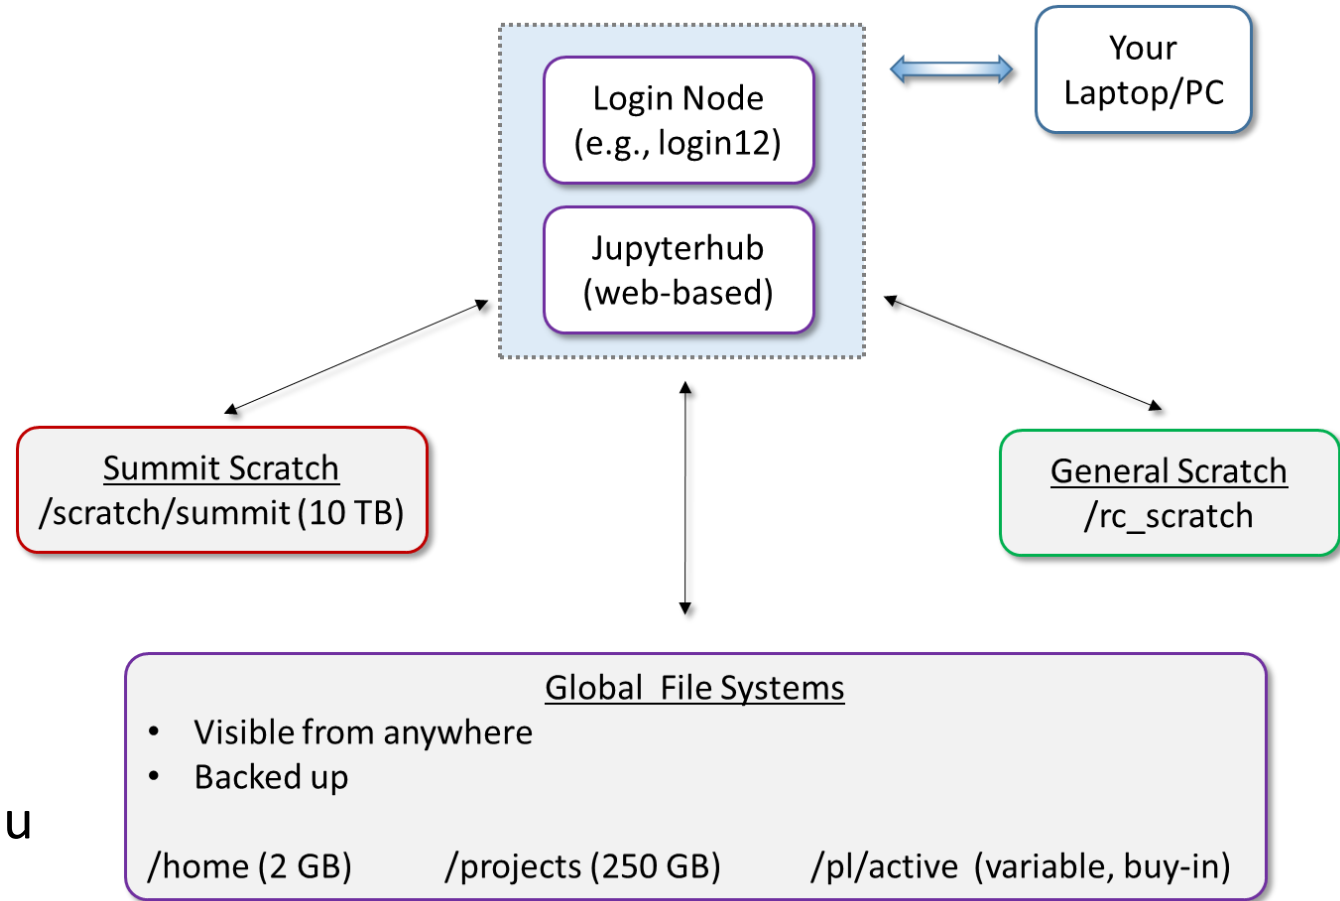
- *Temporary* storage
- Use for program output
- Parallel
- Purged
 - every 90 days
 - per file, based on creation date
- Purged files cannot be retrieved!
- Summit scratch
 - /scratch/summit/username
 - 10 TB
- Blanca scratch
 - /rc_scratch/username
 - Variable storage

RC Systems Conceptual Schematic



Logging in

- Login via ssh with DUO two-factor
 - ssh username@login.rc.Colorado.edu
 - Check your phone for Duo!
- Let's poke around for a bit...
 - cd filesystem/username
 - e.g. cd /scratch/summit/feathern
- Logout (type logout or exit)
- Login in as tutorial user:
 - ssh username@tlogin1.rc.colorado.edu
 - no DUO two-factor



Clone the Seminar Respository

- All of today's materials, including these slides are available online at:
 - https://github.com/feathern/rc_appm_2019
- Let's work in /scratch/summit:
 - ssh scompile (puts you on a compile node)
 - cd /scratch/summit/username
 - git clone https://github.com/feathern/rc_appm_2019.git
 - cd rc_appm_2019

The Software Stack

- Software controlled through a hierarchical module system
- <https://curc.readthedocs.io/en/latest/compute/modules.html>
- Must load a software module to use the module
 - PATH, LD_LIBRARY_PATH etc. modified (for those in the know)
- Many versions of same software may exist, but compiled with
 - Different compiler flavors and versions
(e.g., Intel 17.02 or Intel 16.1 or GNU 7.1)
 - Different MPI flavors and versions
(e.g., Intel MPI or OpenMPI)
- For this reason, modules must be loaded in order:
 1. Compiler
 2. MPI
 3. Desired software
- Let's have a look...

Loading Software Modules

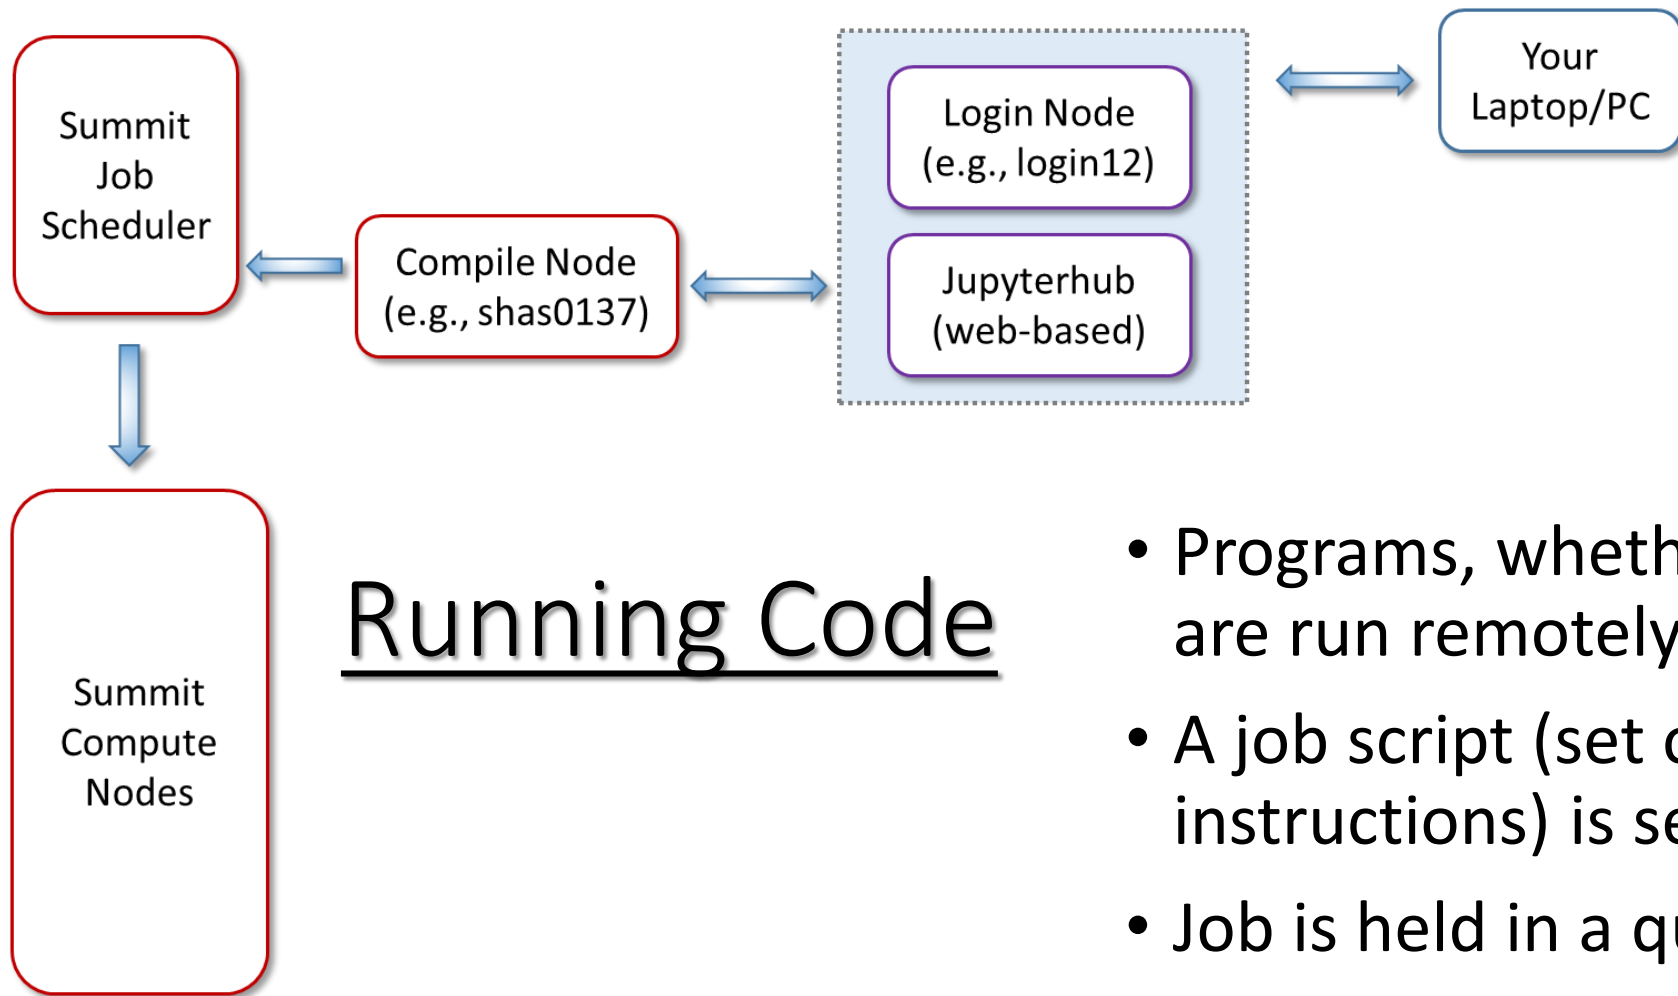
- Primary commands (try these out)
 - module avail
 - See which modules are now available
 - Based on compiler and MPI loaded
 - As you load a compiler and MPI, new modules appear as available
 - module load [or unload] {module name} (e.g., module load intel)
 - Load or unload a desired software module
 - module list
 - See all modules you currently have loaded
 - module purge
 - Unload all currently loaded modules
 - module spider {module name}
 - Search for software (shows which modules need to be loaded)

Compiling Code

- Programs can be compiled within a job script (more on that soon)
- Programs may also be compiled on the scompile nodes
 - You are already on an scompile node
 - Reached from login node via “ssh scompile” (shas prompt shows)
- When compiling a program written in C or Fortran:
 - You must first load the desired compiler/MPI modules
 - These same modules must also be loaded in your *job script*
- Let's try a quick example

Compiling Code

- Let's compile a sample program.
- `cd sample_code`
- Load a compiler and MPI:
 - `module purge`
 - `module load intel`
 - `module load impi`
- Compile the program:
 - `mpicc main.c -o hello.out`
- To run the program, we need to interface with the job scheduler...



Running Code

- Programs, whether parallel or serial, are run remotely
- A job script (set of step-by-step instructions) is sent to the scheduler
- Job is held in a queue
- When enough cores available, it runs...
- Priority basis, sometimes you wait...
- Just like all national level HPC systems...

Job Scripts

- Set of step-by-step instructions necessary to run a job
 - Just like you would type at the command line yourself
- `cd ../sample_scripts`
- `more tutorial_2core.sh`
- Equivalent Summit and Blanca scripts are provided

Job Scripts: Important Notes

- Total number of cores is controlled by *ntasks* flag
- Be sure to specify enough nodes via the *nodes* flag (24 core/node).
- Never request more cores/nodes than you need (node can be shared).
- ALWAYS write to /scratch/summit or /rc_scratch (parallel filesystems)
- NEVER write to /home or /projects from a remote job
 - These systems are not parallel; can negatively impact other users; account locked
- Instructions are executed from within directory where script is submitted
 - So be sure to run appropriate “cd” commands
- Don't forget to load necessary modules!

Running a Remote Job

- Submit the two sample jobs:
 - sbatch tutorial_2core.sh
 - sbatch tutorial_8core.sh
- Job monitoring:
 - squeue -u username (check status of all of your jobs)
- If you goof:
 - scancel jobID# (cancel a job)
- Output normally written to the screen will appear in output file.
- Data output location depends on application

JupyterHub

- Web-based interface to RC Resources
 - Primarily intended for Jupyter notebooks
 - Terminal sessions also available
 - Jupyter sessions spawned on either Summit or Blanca
 - Useful for small quick file transfers (no need to bring ssh into equation)
- Documentation:
<https://curc.readthedocs.io/en/latest/gateways/jupyterhub.html>
- Let's check it out:
 - Go here: <https://jupyter.rc.colorado.edu>
 - Login with YOUR credentials; Click "Start my Server;" Wait.

Globus

- We don't quite have time to discuss this
- If you have questions, feel free to ask me or email rc-help@colorado.edu
- Very useful tool for transferring data between RC and:
 - Your laptop/desktop
 - Another remote HPC system (initiate on your laptop, then walk away)
- <https://curc.readthedocs.io/en/latest/compute/data-transfer.html>